

# COMPSCI 689

## Lecture 18: Autoencoders

Benjamin M. Marlin

College of Information and Computer Sciences  
University of Massachusetts Amherst

Slides by Benjamin M. Marlin (marlin@cs.umass.edu).


# Latent Linear Models

- Latent linear models like factor analysis provide probability densities for linear manifolds in high dimensional spaces.
- Their main limitation is that they can only model linear manifolds.
- Generalized latent linear models can use different latent and visible distributions, but they are still inherently linear.
- To build models for non-linear manifolds, we can apply basis expansions or look at kernelized latent linear models.
- However, these approaches require us to know the right expansion/kernel to apply.

# Autoencoders

- An autoencoder is a deterministic model that consists of two components: an encoder function and a decoder function.
- The encoder function  $f(\mathbf{x})$  maps a  $D$ -dimensional input vector  $\mathbf{x} \in \mathcal{X}$  into a  $K$ -dimensional code vector  $\mathbf{h} \in \mathcal{H}^K$ .
- The decoder function maps the  $K$ -dimensional code vector  $\mathbf{h} \in \mathcal{H}^K$  back to a  $D$ -dimensional reconstruction of the input vector  $\mathbf{r} \in \mathcal{X}$ .
- The goal is typically to learn to accurately copy  $\mathbf{x}$  to  $\mathbf{r}$  while constraining  $\mathbf{h}$  in some way that forces it to encode salient features of the input while ignoring noise.

# Example: Autoencoder for Digits



# Linear Autoencoders

- A linear auto-encoder is an autencoder where the encoder and decoder are linear functions.
- The encoding and decoding functions are:


$$f(\mathbf{x}) = \mathbf{V}\mathbf{x}$$

$$g(\mathbf{h}) = \mathbf{W}\mathbf{h}$$

- Such a linear encoder-decoder model can be learned by minimizing the MSE between the inputs and the reconstructions, often called the *reconstruction error*.

$$\mathbf{V}^*, \mathbf{W}^* = \arg \min_{\mathbf{V}, \mathbf{W}} \sum_{n=1}^N \|\mathbf{x}_n - g(f(\mathbf{x}_n))\|^2$$

# Example: Basic Linear Autoencoder Model



# Factor Analysis as an Autoencoder

- The factor analysis model can be seen as a particularly complex parameterization of a linear autoencoder trained using maximum likelihood estimation (the equations below assume  $\mu = 0$ )
- The encoder function is simply the mean of  $P(\mathbf{z}|\mathbf{x})$ :

$$f(\mathbf{x}) = \mathbb{E}_{P(\mathbf{z}|\mathbf{x})}[\mathbf{z}] = \mathbf{V}\mathbf{x}$$
$$\mathbf{V} = (I + \mathbf{W}^T \Psi \mathbf{W})^{-1} \mathbf{W}^T \Psi^{-1}$$

- The decoder function is the mean of  $P(\mathbf{x}|\mathbf{z})$ :

$$g(\mathbf{z}) = \mathbb{E}_{P(\mathbf{x}|\mathbf{z})}[\mathbf{x}] = \mathbf{W}\mathbf{z}$$

# PPCA as an Autoencoder

- The Probabilistic Principal Components Analysis (PPCA) model is a special case of Factor Analysis where  $\Psi = \sigma^2 I$  and  $\mathbf{W}$  is constrained to be an orthonormal matrix.
- The encoder function is the mean of  $P(\mathbf{z}|\mathbf{x})$ , but the isotropic assumption results in simplifications:

$$f(\mathbf{x}) = \mathbb{E}_{P(\mathbf{z}|\mathbf{x})}[\mathbf{z}] = \mathbf{V}\mathbf{x}$$
$$\mathbf{V} = (\sigma^2 I + \mathbf{W}^T \mathbf{W})^{-1} \mathbf{W}^T$$

- The decoder function is simply the mean function of  $P(\mathbf{x}|\mathbf{z})$ :

$$g(\mathbf{z}) = \mathbb{E}_{P(\mathbf{x}|\mathbf{z})}[\mathbf{x}] = \mathbf{W}\mathbf{z}$$



# Classical Principal Components Analysis as an Autoencoder

- In the limit as  $\sigma^2$  goes to zero, we obtain the classical PCA model. The encoder further simplifies due to orthogonality of  $\mathbf{W}$ :

$$\begin{aligned}f(\mathbf{x}) &= \mathbb{E}_{P(\mathbf{z}|\mathbf{x})}[\mathbf{z}] = \mathbf{V}\mathbf{x} \\ \mathbf{V} &= (\sigma^2 I + \mathbf{W}^T \mathbf{W})^{-1} \mathbf{W}^T \rightarrow \mathbf{W}^T\end{aligned}$$

- The decoder function is again the mean of  $P(\mathbf{x}|\mathbf{z})$ :

$$g(\mathbf{z}) = \mathbb{E}_{P(\mathbf{x}|\mathbf{z})}[\mathbf{x}] = \mathbf{W}\mathbf{z}$$

- This result shows that classical PCA is actually a directly parameterized linear autoencoder. General linear autoencoders do not require orthogonality of  $\mathbf{W}$  and  $\mathbf{V}$ , but they identify exactly the same linear subspace!

# Classical Principal Components Analysis

- One of the interesting properties of classical PCA is that you do not need iterative numerical optimization to find the optimal  $\mathbf{W}$ .
- Let  $\mathbf{X}$  be an  $N \times D$  data matrix that has been mean centered.
- Define the empirical scatter matrix to be  $\mathbf{S} = \mathbf{X}^T \mathbf{X}$ .
- Then the optimal  $\mathbf{W}$  is given by the leading  $K$  eigenvectors of  $\mathbf{S}$ .
- These eigenvectors are orthogonal by definition and are the rank  $K$  maximum variance sub-space of  $\mathbf{X}$ , also called the *principal sub-space*.
- Interestingly, this estimation approach also identifies the optimal  $\mathbf{W}$  for the PPCA model.
- Finally, the rank  $K$  singular value decomposition of  $\mathbf{X}$  is given by  $\mathbf{USV}^T$  where  $\mathbf{U}$  is an  $N \times K$  orthonormal matrix,  $\mathbf{U}$  is a  $D \times K$  orthonormal matrix and  $\mathbf{S}$  is a diagonal matrix.

# Singular Value Decomposition

- The rank  $K$  singular value decomposition (SVD) of  $\mathbf{X}$  is given by  $\mathbf{USV}^T$  where  $\mathbf{U}$  is an  $N \times K$  orthonormal matrix,  $\mathbf{U}$  is a  $D \times K$  orthonormal matrix and  $\mathbf{S}$  is a positive diagonal matrix.
- The rank  $K$  SVD of  $\mathbf{X}$  minimizes the objective function  $\|\mathbf{X} - \mathbf{USV}^T\|_2^2$  subject to the stated conditions on  $\mathbf{U}, \mathbf{S}, \mathbf{V}$ .
- The matrix  $\mathbf{V}^T$  identified using the SVD is exactly equal to the  $\mathbf{W}$  matrix identified using PCA (or PPCA).
- This means that the SVD provides yet another way of identifying the optimal parameters for a linear autoencoder.

# Constraints on Linear Autoencoders

- In order for an MSE-minimizing linear autoencoder to learn something useful about the structure of the input data, it is necessary to constrain the code vector in some way.
- If  $K = D$  (complete representation) or  $K > D$  (overcomplete representation), then we can achieve zero reconstruction error by setting the first  $D$  rows of  $V$  to the identity matrix and the first  $K$  columns of  $W$  to the identity matrix (all other entries are zero).
- One way to constrain a linear autoencoder is to require  $K < D$  (an undercomplete representation).
- This dimensionality reduction forces the model to extract useful information from the inputs and to discard noise in order to minimize the reconstruction error.

# Sparse Overcomplete Autoencoders

- A different way to constrain a linear autoencoder is to allow  $K > D$  while constraining the number of non-zero elements of the code vector  $\mathbf{h}$ .
- *Sparse coding* is a dimensionality reduction model that has a linear decoder and an optimization-based encoder:

$$f(\mathbf{x}) = \arg \min_{\mathbf{h}} \|\mathbf{x} - \mathbf{W}\mathbf{h}\|_2^2 + \lambda \|\mathbf{h}\|_1$$

$$g(\mathbf{h}) = \mathbf{W}\mathbf{h}$$

- The parameters  $\mathbf{W}$  are learned on a data set using:

$$\arg \min_{\mathbf{W}, \mathbf{h}_1, \dots, \mathbf{h}_N} \sum_{n=1}^N (\|\mathbf{x}_n - \mathbf{W}\mathbf{h}_n\|_2^2 + \lambda \|\mathbf{h}_n\|_1)$$

# Sparse Overcomplete Linear Autoencoders

- Sparse coding can achieve sparse code vectors through direct  $\ell_1$  penalization during encoding, but they need to perform optimization to infer the code can be prohibitive.
- An alternative is to learn a linear autoencoder while penalizing the  $\ell_1$  norm of the codes it produces to constrain the complexity of the codes the model tends to generate.
- The resulting optimization problem looks like:

$$\mathbf{V}^*, \mathbf{W}^* = \arg \min_{\mathbf{V}, \mathbf{W}} \sum_{n=1}^N (\|\mathbf{x}_n - g(f(\mathbf{x}_n))\|_2^2 + \lambda \|f(\mathbf{x}_n)\|_1)$$

# Linear Denoising Autoencoders

- Yet another way to stop an overcomplete autoencoder from simply learning the identity function is to make the problem that the autoencoder has to solve more difficult.
- A *denoising autoencoder* purposely corrupts the input  $\mathbf{x}$  using a sample drawn from a stochastic noise process  $q(\mathbf{x}'|\mathbf{x})$ .
- It then provides  $\mathbf{x}'$  as the input to the autoencoder while requiring the reconstruction that the model produced match the original  $\mathbf{x}$ .
- The model can be trained as shown below. The characteristics of the noise process are hyper-parameters of model.


$$\mathbf{V}^*, \mathbf{W}^* = \arg \min_{\mathbf{V}, \mathbf{W}} \sum_{n=1}^N \mathbb{E}_{q(\mathbf{x}'|\mathbf{x}_n)} [\|\mathbf{x}_n - g(f(\mathbf{x}'))\|_2^2]$$

# Non-Linear Autoencoders

- All of the above models and training criteria can only accurately represent data defined on linear manifolds.
- To make an autoencoder non-linear, it suffices to make the encoder and decoder networks non-linear.
- However, the capacity of non-linear autoencoders is no longer limited by the length of the code vector.
- In fact, it is possible to construct a deep network where the code is 1-dimensional, but the network achieves zero reconstruction loss.



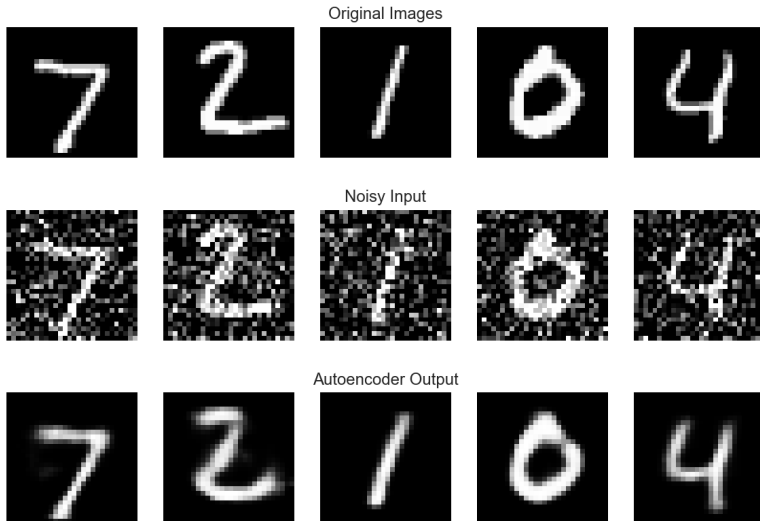
# Example: Deep Autoencoders



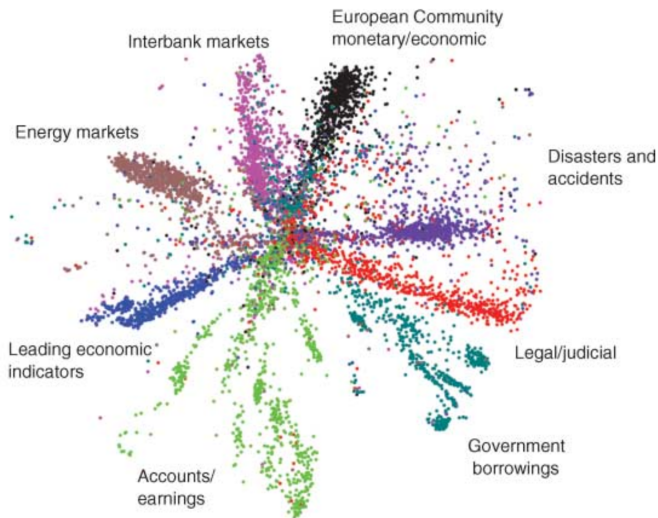
# Constraining Non-Linear Autoencoders

- Like with overcomplete linear autoencoders, nonlinear autoencoders need to have constraints to force them to extract useful information from the data.
- Possible constraints include constraining the depth of the encoder and decoder networks as well as the length of the code vectors, using sparse codes, and using the denoising principle.
- As with supervised learning, it is known that some data distributions can be much more efficiently represented with deep, non-linear autoencoders instead of shallow non-linear or linear autoencoders.
- The key is determining the network architecture that gives the best performance for a given down-stream task.

# Example: Deep Denoising Autoencoder for Images



# Example: Deep Autoencoder Embedding for Documents



# Autoencoder Applications

- Autoencoders have many applications including feature extraction, image and signal denoising, and semantic hashing for information retrieval.
- Unsupervised feature extraction (e.g., representation learning) is one of their most common applications and involves learning an encoder-decoder pair on a large unlabeled data set.
- The basic autoencoder architecture can also be modified for semi-supervised learning by including an encoder, a decoder, and a classifier that uses the code vector as its input.

# Example: Semi-Supervised Autoencoder/Classifier

