

# COMPSCI 689

## Lecture 19: Generalizing Autoencoders and RBMs

Benjamin M. Marlin

College of Information and Computer Sciences  
University of Massachusetts Amherst

Slides by Benjamin M. Marlin (marlin@cs.umass.edu).



# Factor Analysis to Autoencoders

- Last class, we saw how non-linear autoencoder neural network models generalize latent linear models (i.e., Factor Analysis).
- This allows autoencoders to model real-valued data defined on non-linear real manifolds.
- We also previously saw how factor analysis models can be generalized to deal with other inputs.
- In this class we'll provide a similar generalization for autoencoders, and introduce restricted Boltzmann machines.

# Generalized Autoencoders


- We can use the same ideas we applied to generalize factor analysis to generalize autoencoders to different data types.
- In general, the input to first hidden layer structure remains the same (an exception is categorical variables, which need to be converted to a one-hot encoding).
- The hidden-to-output structure can be interpreted as computing the parameters of an exponential family distribution for each  $X_d$  using the last layer of hidden values  $\mathbf{h}^L = [h_1^L, \dots, h_{K_L}^L]$ .

## Example: Binary Autoencoders

- Suppose  $\mathbf{X} = [X_1, \dots, X_D]$  consists of binary values only.
- Let the last layer of hidden values be  $\mathbf{h}^L = [h_1^L, \dots, h_{K_L}^L]$ .
- Then we can define the probability of the reconstruction  $\mathbf{r}$  given the last layer of hidden unit values  $\mathbf{h}^L$  as follows:

$$\begin{aligned}\theta_{d|h} &= \frac{1}{1 + \exp(-\mathbf{W}_d^L \mathbf{h}^L)} \\ P(r_d | \mathbf{h}^L) &= \theta_{d|h}^{r_d} (1 - \theta_{d|h})^{(1-r_d)} \\ P(\mathbf{R} = \mathbf{r} | \mathbf{h}^L) &= \prod_{d=1}^D P(r_d | \mathbf{h}^L)\end{aligned}$$

# Example: Deep Generalized Autoencoders



# Learning Deep Generalized Autoencoders

- To learn a deep generalized autoencoder, we optimize the objective function shown below where we recall that  $\mathbf{h}_n^L$  is computed using a forward pass through the autoencoder starting from input  $\mathbf{x}_n$ :

$$\mathcal{J}(\mathcal{D}, \mathbf{W}^{1:L}) = \sum_{n=1}^N \log P(\mathbf{R} = \mathbf{x}_n | \mathbf{h}_n^L)$$

- This function effectively replaces the mean squared error between the input and the output used in a standard autoencoder with a log loss derived from  $P(\mathbf{R} = \mathbf{x}_n | \mathbf{h}_n^L)$ .
- Using a Gaussian distribution will recover squared loss. Using a Laplacian distribution will recover absolute loss. A product of Bernoulli's will result in cross-entropy loss.
- However, the resulting code in the middle of the autoencoder

# Learning Deep Generalized Autoencoder Codes

- However, there are no latent variables to integrate or optimize over in a deep generalized autoencoder. The code vector is a deterministic function of the input.
- However, the resulting code in the middle of the a deep generalized autoencoder is still real-valued.
- Unlike in the generalized factor analysis model family, there is no way to obtain other kinds of codes.

# Restricted Boltzmann Machines

- RBMs are closely related to the factor analysis model family.
- They are also closely related to a class of models from statistical physics called Ising models or spin glass models that are joint models for binary variables.
- They are also closely related to a class of models studied in statistics called Markov Random Fields, which generalize Ising models and Markov chains.
- In the ML literature, these models are described using connectionist terminology.
- In the classical RBM,  $\mathbf{h} \in \{0, 1\}^K$  is a vector of hidden units while  $\mathbf{x} \in \{0, 1\}^D$  is a vector of observed values. The model parameterizes  $P(\mathbf{X} = \mathbf{x}, \mathbf{H} = \mathbf{h})$ .



# Restricted Boltzmann Machines

- Instead of being composed of two locally normalized probability distributions (e.g.,  $P(X|Z)$  and  $P(Z)$ ), they are constructed from a joint *energy function* that requires explicit normalization. These models are also called energy-based models (EBMs).

$$E_W^k(\mathbf{x}, h_k) = - \sum_{d=1}^D (W_{dk}^P x_d h_k + W_k^B h_k)$$

$$E_W(\mathbf{x}, \mathbf{h}) = \sum_{k=1}^K E_W^k(\mathbf{x}, h_k)$$

$$P_W(\mathbf{X} = \mathbf{x}, \mathbf{H} = \mathbf{h}) = \frac{\exp(-E_W(\mathbf{x}, \mathbf{h}))}{\sum_{\mathbf{x}'} \sum_{\mathbf{h}'} \exp(-E_W(\mathbf{x}', \mathbf{h}'))}$$


# Restricted Boltzmann Machines

- We can equivalently write the model as a product of distributions, one for each hidden unit. For this reason, the model is sometimes also called a product of experts (PoE).


$$P_W(\mathbf{X} = \mathbf{x}, \mathbf{H} = \mathbf{h}) = \frac{\prod_{k=1}^K \exp(-E_W^k(\mathbf{x}, h_k))}{\sum_{\mathbf{x}'} \sum_{\mathbf{h}'} \prod_{k'=1}^K \exp(-E_W^{k'}(\mathbf{x}', h_{k'}))}$$

# Products vs Mixtures

## A picture of the two combination methods



Mixture model:  
Scale each  
distribution down  
and add them  
together



Product model:  
Multiply the two  
densities together  
at every point and  
then renormalize.

Slide from Geoff Hinton, CSC321, Lecture 24

# Restricted Boltzmann Machines: Conditionals

Unlike for the factor analysis family, both  $P(\mathbf{X}|\mathbf{H} = \mathbf{h})$  and  $P(\mathbf{H}|\mathbf{X} = \mathbf{x})$  are fully factorized distributions that are easily computable:

$$P_W(\mathbf{X} = \mathbf{x}|\mathbf{h}) = \prod_{d=1}^D P_W(X_d = x_d|\mathbf{h})$$

$$P_W(\mathbf{H} = \mathbf{h}|\mathbf{x}) = \prod_{k=1}^K P_W(H_k = h_k|\mathbf{x})$$

# Restricted Boltzmann Machines: Conditionals

Unlike for the factor analysis family, both  $P(\mathbf{X}|\mathbf{H} = \mathbf{h})$  and  $P(\mathbf{H}|\mathbf{X} = \mathbf{x})$  are fully factorized distributions that are easily computable:

$$P_W(X_d = x_d|\mathbf{h}) = \frac{\exp\left(\sum_{k=1}^K W_{dk}^P x_d h_k\right)}{1 + \exp\left(\sum_{k=1}^K W_{dk}^P h_k\right)}$$
$$P_W(H_k = h_k|\mathbf{x}) = \frac{\exp\left(\sum_{d=1}^D W_{dk}^P x_d h_k + W_k^B h_k\right)}{1 + \exp\left(\sum_{d=1}^D W_{dk}^P x_d + W_k^B\right)}$$

# Restricted Boltzmann Machines: Visible Marginal

- It's also actually possible to sum over all joint settings of the hidden variables analytically to get an expression for  $P_W(\mathbf{X})$ , but these values are typically not computable:


$$P_W(\mathbf{X} = \mathbf{x}) = \frac{\prod_{k=1}^K \left( 1 + \exp \left( \sum_{d=1}^D W_{dk}^P x_d + W_k^B \right) \right)}{\sum_{\mathbf{x}'} \prod_{k=1}^K \left( 1 + \exp \left( \sum_{d=1}^D W_{dk}^P x'_d + W_k^B \right) \right)}$$

# Restricted Boltzmann Machines: Stochastic MLE

- Again, unlike the FA model family, the RBM family can be learned using marginal likelihood maximization, but sampling is required to efficiently approximate marginal likelihood gradients. The result is a stochastic maximum marginal likelihood algorithm.

$$\frac{\partial \mathcal{L}(\mathcal{D}, \mathbf{W})}{\partial W_{dk}^P} = \frac{1}{N} \sum_{n=1}^N x_{dn} P(H_k = 1 | \mathbf{x}_n) - \frac{1}{S} \sum_{s=1}^S \hat{x}_{ds} P(H_k = 1 | \hat{\mathbf{x}}_s)$$
$$\hat{\mathbf{x}}_s \sim P_W(\mathbf{X})$$

- The sampling procedure alternates between sampling  $\hat{\mathbf{x}}$  and  $\hat{\mathbf{h}}$  using the structure shown below. Only the  $\hat{\mathbf{x}}$  samples are used for learning.





# Restricted Boltzmann Machines: Applications

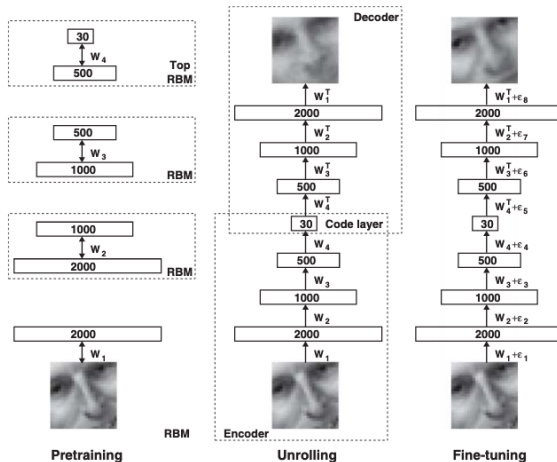


Image from Hinton and Salakhutdinov. *Reducing the Dimensionality of Data with Neural Networks*. Science. 2006.